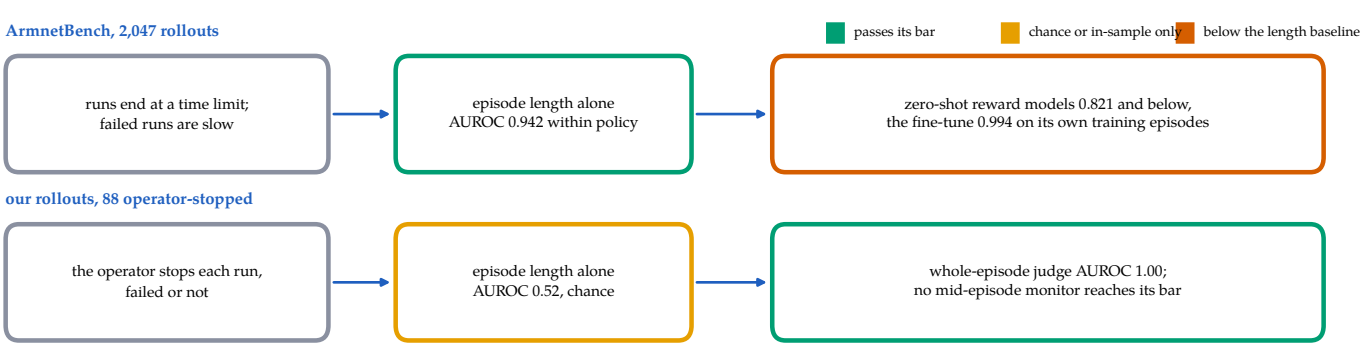
ArmnetBench, 2,047 rollouts
passes its bar
chance or in-sample only
below the length baseline
runs end at a time limit;
failed runs are slow
episode length alone
AUROC 0.942 within policy
zero-shot reward models 0.821 and below,
the fine-tune 0.994 on its own training episodes
our rollouts, 88 operator-stopped
the operator stops each run,
failed or not
episode length alone
AUROC 0.52, chance
whole-episode judge AUROC 1.00;
no mid-episode monitor reaches its bar